

Intro to Analytics Modeling HW 6

2024-06-26

Question 13.2

Please note that this will be available at the end of the document. I used SimPy in Python and needed to add it after I've completed the other questions in R.

Importing Libraries

```
library(DAAG)
```

Read Data

```
breast_cancer_data <- read.table("~/Desktop/ISYE-6501/week 6 Homework-Summer24/breast-cancer-wisconsin.data.txt")
breast_cancer_df <- data.frame(breast_cancer_data)
# convert "?" to NA
breast_cancer_df[breast_cancer_df == "?"] <- NA
head(breast_cancer_df)
```

```
##      X1000025 X5 X1 X1.1 X1.2 X2 X1.3 X3 X1.4 X1.5 X2.1
## 1    1002945  5  4     4     5  7   10  3     2     1     2
## 2    1015425  3  1     1     1  2    2  3     1     1     2
## 3    1016277  6  8     8     1  3    4  3     7     1     2
## 4    1017023  4  1     1     3  2    1  3     1     1     2
## 5    1017122  8 10    10     8  7   10  9     7     1     4
## 6    1018099  1  1     1     1  2   10  3     1     1     2
```

Data Cleansing

```
breast_cancer_df$X1.3 <- as.numeric(breast_cancer_df$X1.3)
```

Question 14.1

The breast cancer data set `breast-cancer-wisconsin.data.txt` from <http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/> (description at <http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Original%29>) has missing values.


check each column if there's any missing values

```
cols <- colnames(breast_cancer_df)[2:11]
for (col in cols) {
  if (any(is.na(breast_cancer_df[[col]]))) {
    print(paste(col, "has at least one missing value"))
  }
}
```

```
## [1] "X1.3 has at least one missing value"
```


1. Use the mean/mode imputation method to impute values for the missing data.

```
# Create a histogram
hist(breast_cancer_df$X1.3, main = "Histogram of Data", xlab = "Values", ylab = "Frequency", col = "lightblue")
```



```
qqnorm(breast_cancer_df$X1.3)
qqline(breast_cancer_df$X1.3, col = "red")
```

Normal Q-Q Plot



```
unique(breast_cancer_df$X1.3)
```

```
## [1] 10 2 4 1 3 9 7 NA 5 8 6
```

Observation: To determine which statistic would be better between mean and mode, I plotted a histogram, a Q-Q plot, and checked the data values to determine its type (numerical vs. categorical) and assess its normality in case it's numerical data. Based on the visualization outcomes and checking unique values, it appears that the column represents categorical data, making mode a more appropriate choice.

```
# Get the mode and replace NA with the mode
get_mode <- function(x) {
  ux <- unique(x)
  ux[which.max(tabulate(match(x, ux)))]
}

breast_cancer_df_mode <- breast_cancer_df
breast_cancer_df_mode$X1.3 <- as.integer(breast_cancer_df_mode$X1.3)

x1.3_mode <- get_mode(breast_cancer_df_mode$X1.3)
print(x1.3_mode)
```

```
## [1] 1
```

```
# replace NA with the mode (=1)
breast_cancer_df_mode$X1.3[is.na(breast_cancer_df_mode$X1.3)] <- x1.3_mode
```

```
# Check if imputation using mode (=1) works fine.
cols <- colnames(breast_cancer_df_mode)[2:11]
for (col in cols) {
  if (any(is.na(breast_cancer_df_mode[[col]]))) {
    print(paste(col, "has at least one missing value"))
  } else {
    print(paste(col, "has no missing values"))
  }
}
```

```
## [1] "X5 has no missing values"
## [1] "X1 has no missing values"
## [1] "X1.1 has no missing values"
## [1] "X1.2 has no missing values"
## [1] "X2 has no missing values"
## [1] "X1.3 has no missing values"
## [1] "X3 has no missing values"
## [1] "X1.4 has no missing values"
## [1] "X1.5 has no missing values"
## [1] "X2.1 has no missing values"
```

```
# (Optional) Get the mean and replace NA with the mode if desired.
breast_cancer_df_mean <- breast_cancer_df
breast_cancer_df_mean$X1.3 <- as.numeric(breast_cancer_df_mean$X1.3)
x1.3_mean <- mean(breast_cancer_df_mean$X1.3, na.rm = TRUE)
print(x1.3_mean)
```

```
## [1] 3.548387
```

```
# replace NA with the mean
breast_cancer_df_mean$X1.3[is.na(breast_cancer_df_mean$X1.3)] <- x1.3_mean
head(breast_cancer_df_mean)
```

```
##      X1000025 X5 X1 X1.1 X1.2 X2 X1.3 X3 X1.4 X1.5 X2.1
## 1   1002945  5  4    4    5  7   10  3    2    1    2
## 2   1015425  3  1    1    1  2    2  3    1    1    2
## 3   1016277  6  8    8    1  3    4  3    7    1    2
## 4   1017023  4  1    1    3  2    1  3    1    1    2
## 5   1017122  8 10   10    8  7   10  9    7    1    4
## 6   1018099  1  1    1    1  2   10  3    1    1    2
```

```
# Check if imputation using mode (=1) works fine.
cols <- colnames(breast_cancer_df_mean)[2:11]
for (col in cols) {
  if (any(is.na(breast_cancer_df_mean[[col]]))) {
    print(paste(col, "has at least one missing value"))
  } else {
    print(paste(col, "has no missing values"))
  }
}
```

```
## [1] "X5 has no missing values"
## [1] "X1 has no missing values"
## [1] "X1.1 has no missing values"
## [1] "X1.2 has no missing values"
## [1] "X2 has no missing values"
## [1] "X1.3 has no missing values"
## [1] "X3 has no missing values"
## [1] "X1.4 has no missing values"
## [1] "X1.5 has no missing values"
## [1] "X2.1 has no missing values"
```

2. Use regression to impute values for the missing data.

```
na_indices <- which(is.na(breast_cancer_df$X1.3))
breast_cancer_df_cleansed <- breast_cancer_df[-na_indices,]
lm_model <- lm(X1.3~X5+X1+X1.1+X1.2+X2+X3+X1.4+X1.5+X2.1, data=breast_cancer_df_cleansed)
summary(lm_model)
```

```
##
## Call:
## lm(formula = X1.3 ~ X5 + X1 + X1.1 + X1.2 + X2 + X3 + X1.4 +
##      X1.5 + X2.1, data = breast_cancer_df_cleansed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.6037 -0.4280 -0.2197  0.8720  8.6286
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.25090    0.31013  -13.707  < 2e-16 ***
## X5              0.01890    0.03968   0.476  0.63407
## X1             -0.16224    0.06736  -2.408  0.01629 *
## X1.1            0.18420    0.06556   2.810  0.00510 **
## X1.2            0.22088    0.04128   5.351  1.2e-07 ***
## X2              0.01924    0.05527   0.348  0.72785
## X3              0.15163    0.05336   2.842  0.00462 **
## X1.4           -0.08743    0.03972  -2.201  0.02807 *
## X1.5           -0.06300    0.05222  -1.206  0.22810
## X2.1            2.50879    0.17831  14.070  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.002 on 672 degrees of freedom
## Multiple R-squared:  0.7025, Adjusted R-squared:  0.6985
## F-statistic: 176.3 on 9 and 672 DF, p-value: < 2.2e-16
```

```
# Rebuild a linear regression model using significant predictors based on the p-value
lm_model <- lm(X1.3~X1+X1.1+X1.2+X3+X1.4+X2.1, data=breast_cancer_df_cleansed)
summary(lm_model)
```


```
##
```

```
## Call:
## lm(formula = X1.3 ~ X1 + X1.1 + X1.2 + X3 + X1.4 + X2.1, data = breast_cancer_df_cleansed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.9671 -0.4143 -0.2570  0.8576  8.5857
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.28598    0.30453  -14.074 < 2e-16 ***
## X1            -0.16305    0.06504   -2.507  0.01241 *
## X1.1           0.18760    0.06502    2.885  0.00403 **
## X1.2           0.21361    0.04071    5.247 2.07e-07 ***
## X3             0.15728    0.05307    2.964  0.00315 **
## X1.4          -0.09287    0.03910   -2.375  0.01782 *
## X2.1           2.54159    0.16404   15.494 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.999 on 675 degrees of freedom
## Multiple R-squared:  0.7018, Adjusted R-squared:  0.6991
## F-statistic: 264.7 on 6 and 675 DF, p-value: < 2.2e-16

# Use cross-validation to test how good this model really is.
model_cv <- cv.lm(breast_cancer_df_cleansed, lm_model, m=5, printit = FALSE)
```

```
## Warning in cv.lm(breast_cancer_df_cleansed, lm_model, m = 5, printit = FALSE):
##
## As there is >1 explanatory variable, cross-validation
## predicted values for a fold are not a linear function
## of corresponding overall predicted values. Lines that
## are shown for the different folds are approximate
```

Small symbols show cross-validation predicted values



```
SST <- sum(
  (
    as.numeric(breast_cancer_df_cleansed$X1.3) - mean(as.numeric(breast_cancer_df_cleansed$X1.3)))^2
  )
R2_cv <- 1 - attr(model_cv, "ms") * nrow(breast_cancer_df) / SST
print(R2_cv)
```

```
## [1] 0.6827687
```

Observation: R squared value from the cross validation looks pretty good, which means that the linear fit would work pretty well to predict the missing values.

```
# Get predictions for missing X1.3 values
X1.3_hat <- predict(lm_model, newdata = breast_cancer_df[na_indices,])
print(X1.3_hat)
```

```
##      23      40     139     145     158     164     235     249
## 7.2021358 3.2249842 1.2570453 1.6322455 1.2873664 1.4143244 1.9771246 1.4143244
##      275     292     294     297     315     321     411     617
## 1.6322455 6.3053734 1.2570453 0.9573642 1.8318484 1.4143244 1.2570453 1.0997663
```

```
print(round(X1.3_hat))
```

```
##  23  40 139 145 158 164 235 249 275 292 294 297 315 321 411 617
##   7   3   1   2   1   1   2   1   2   6   1   1   2   1   1   1
```

```

# Impute X1.3 for missing values using predicted values from the linear regression model
breast_cancer_df_reg_imputed <- breast_cancer_df
breast_cancer_df_reg_imputed[na_indices,]$X1.3 <- X1.3_hat
breast_cancer_df_reg_imputed$X1.3 <- as.numeric(breast_cancer_df_reg_imputed$X1.3)

# Round X1.3_hat values since the original values are all integer
breast_cancer_df_reg_imputed[na_indices,]$X1.3 <- round(X1.3_hat)
breast_cancer_df_reg_imputed$X1.3 <- as.integer(breast_cancer_df_reg_imputed$X1.3)

# Make sure no X1.3 values are outside of the given range, which is between 1 and 10
breast_cancer_df_reg_imputed$X1.3[breast_cancer_df_reg_imputed$X1.3 > 10] <- 10
breast_cancer_df_reg_imputed$X1.3[breast_cancer_df_reg_imputed$X1.3 < 1] <- 1
head(breast_cancer_df_reg_imputed)

```

```

##      X1000025 X5 X1 X1.1 X1.2 X2 X1.3 X3 X1.4 X1.5 X2.1
## 1    1002945  5  4     4     5  7    10  3     2     1     2
## 2    1015425  3  1     1     1  2     2  3     1     1     2
## 3    1016277  6  8     8     1  3     4  3     7     1     2
## 4    1017023  4  1     1     3  2     1  3     1     1     2
## 5    1017122  8 10    10     8  7    10  9     7     1     4
## 6    1018099  1  1     1     1  2    10  3     1     1     2

```

3. Use regression with perturbation to impute values for the missing data.

```

# Adding a random amount up or down from the imputed estimate using linear model built above
# This can be done using a random normal distribution. We take a mean and std of the predicted values.
set.seed(1)
X1.3_hat_pert <- rnorm(nrow(breast_cancer_df_reg_imputed[na_indices,]), X1.3_hat, sd(X1.3_hat))
print(X1.3_hat_pert)

```

```

##      [1]  6.0372823  3.5664580 -0.2967562  4.5985750  1.9000664 -0.1112876
##      [7]  2.8834699  2.7871951  2.7028766  5.7375220  4.0681122  1.6822544
##     [13]  0.6766886 -2.7037779  3.3487872  1.0162150

```

```

print(round(X1.3_hat_pert))

```

```

##      [1]  6  4  0  5  2  0  3  3  3  6  4  2  1 -3  3  1

```

```

# Notice that there are some negative values when we perturb the predicted values.
breast_cancer_df_pert_imputed <- breast_cancer_df
breast_cancer_df_pert_imputed[na_indices,]$X1.3 <- round(X1.3_hat_pert)
breast_cancer_df_pert_imputed$X1.3 <- as.integer(breast_cancer_df_pert_imputed$X1.3)

# Make sure no X1.3 values are outside of the given range, which is between 1 and 10
breast_cancer_df_pert_imputed$X1.3[breast_cancer_df_pert_imputed$X1.3 > 10] <- 10
breast_cancer_df_pert_imputed$X1.3[breast_cancer_df_pert_imputed$X1.3 < 1] <- 1
head(breast_cancer_df_pert_imputed)

```

```

##      X1000025 X5 X1 X1.1 X1.2 X2 X1.3 X3 X1.4 X1.5 X2.1

```


## 1	1002945	5	4	4	5	7	10	3	2	1	2
## 2	1015425	3	1	1	1	2	2	3	1	1	2
## 3	1016277	6	8	8	1	3	4	3	7	1	2
## 4	1017023	4	1	1	3	2	1	3	1	1	2
## 5	1017122	8	10	10	8	7	10	9	7	1	4
## 6	1018099	1	1	1	1	2	10	3	1	1	2

Question 15.1

Describe a situation or problem from your job, everyday life, current events, etc., for which optimization would be appropriate. What data would you need?

I work as a data scientist at a company specializing in lithium metal battery manufacturing. One of our primary challenges lies in electrode production, where we face issues with inconsistent quality and high production costs due to low yield. To address this, an optimization initiative is underway to refine process parameters and achieve higher quality electrodes at reduced costs.

The data that I need are as follows:

- process parameters (e.g. electrode dimensions, burr, etc.)
- quality metrics (e.g. HiPot failure rate)
- environmental conditions (e.g. ambient temperature, humidity, etc.)
- machine performance data (e.g. machine up/downtime and preventative maintenance schedule)

Question 13.2

In this problem you, can simulate a simplified airport security system at a busy airport. Passengers arrive according to a Poisson distribution with $\lambda_1 = 5$ per minute (i.e., mean interarrival rate $\mu_1 = 0.2$ minutes) to the ID/boarding-pass check queue, where there are several servers who each have exponential service time with mean rate $\mu_2 = 0.75$ minutes. [Hint: model them as one block that has more than one resource.] After that, the passengers are assigned to the shortest of the several personal-check queues, where they go through the personal scanner (time is uniformly distributed between 0.5 minutes and 1 minute).

Use the Arena software (PC users) or Python with SimPy (PC or Mac users) to build a simulation of the system, and then vary the number of ID/boarding-pass checkers and personal-check queues to determine how many are needed to keep average wait times below 15 minutes. [If you're using SimPy, or if you have access to a non-student version of Arena, you can use $\lambda_1 = 50$ to simulate a busier airport.]

<My Python code for the simulation>

```
1 import simpy
2 import random
3 import numpy as np
4 import pandas as pd
5 import plotly.express as px

6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49

class SecurityCheckpoint:
    def __init__(self, env, id_stations, scan_stations, id_check_server_time, scan_time_range):
        self.env = env
        self.id_check = simpy.Resource(env, id_stations)
        self.scan = simpy.Resource(env, scan_stations)
        self.id_check_server_time = id_check_server_time
        self.scan_time_range = scan_time_range

    def check_id(self, traveler_id):
        id_check_duration = random.expovariate(1 / self.id_check_server_time)
        yield self.env.timeout(id_check_duration)
        # print(f'A traveler {traveler_id} leaves the id check station at {env.now:.2f}.')

    def scan_traveler(self, traveler_id):
        scan_duration = random.uniform(self.scan_time_range[0], self.scan_time_range[1])
        yield self.env.timeout(scan_duration)
        # print(f'A traveler {traveler_id} leaves the scanning station at {env.now:.2f}.')

    def traveler(env, checkpoint, traveler_id):
        arrival_time = env.now

        with checkpoint.id_check.request() as request:
            yield request # a traveler enters id check station
            yield env.process(checkpoint.check_id(traveler_id)) # the traveler leaves id check station

        with checkpoint.scan.request() as request:
            yield request # a traveler enters scanning station
            yield env.process(checkpoint.scan_traveler(traveler_id)) # the traveler leaves id check station
        # print(f'now {env.now} - arrival_time {arrival_time} = {env.now - arrival_time}')
        wait_durations.append(env.now - arrival_time)

def security_checkpoint_operation(env, id_stations, scan_stations, id_check_server_time, scan_time_range):
    checkpoint = SecurityCheckpoint(env, id_stations, scan_stations, id_check_server_time, scan_time_range)

    traveler_id = 1

    # Create 1 initial traveler
    for _ in range(1):
        env.process(traveler(env, checkpoint, f'{traveler_id}'))

    while True:
        yield env.timeout(random.expovariate(1 / MEAN_INTERARRIVAL_RATE))
        traveler_id += 1
        env.process(traveler(env, checkpoint, f'{traveler_id}'))

def average_wait_time(durations):
    mean_wait = np.mean(durations)
    return mean_wait
```

```

1 SEED = 1
2 ID_STATIONS = 20 # will vary
3 ID_CHECK_SERVER_TIME = 0.75 # minutes
4
5 SCAN_STATIONS = 20 # will vary
6 SCAN_TIME_RANGE = [0.5, 1] # minutes
7
8 # Passengers arrive according to a Poisson distribution with  $\lambda 1$  / per minute
9 lambda1 = 50 #  $\lambda 1 = 50$  to simulate a busier airport
10 MEAN_INTERARRIVAL_RATE = 1 / lambda1 # minutes

1 # Getting unique combinations of two lists 'id_check_station_time' and 'scan_station_time'
2 import itertools
3
4 id_check_station_time = list(range(10, 41, 1))
5 scan_station_time = list(range(10, 41, 1))
6 combinations = list(itertools.product(id_check_station_time, scan_station_time))

1 %%time
2 various_configs = list()
3 random.seed(SEED)
4 for comb in combinations:
5     wait_durations = []
6     ID_STATIONS = comb[0]
7     SCAN_STATIONS = comb[1]
8     env = simply.Environment()
9     env.process(security_checkpoint_operation(env, ID_STATIONS, SCAN_STATIONS, ID_CHECK_SERVER_TIME, SCAN_TIME_RANGE))
10    env.run(until=100)
11    avg_waiting = average_wait_time(wait_durations)
12    various_configs.append(
13        {
14            "num_id_stations": ID_STATIONS,
15            "num_scan_stations": SCAN_STATIONS,
16            "avg_waiting": avg_waiting,
17            "config": str(ID_STATIONS) + "&" + str(SCAN_STATIONS)
18        }
19    )
20 del env
21
22 df_various_configs = pd.DataFrame(various_configs)


CPU times: user 1min 31s, sys: 15.6 ms, total: 1min 31s
Wall time: 1min 31s

```

```

1 fig = px.scatter(df_various_configs.sort_values(by=["num_id_stations", "num_scan_stations"]), x="config", y="avg_waiting")
2 fig.add_hline(y=15, line_color="red")
3 fig.show()

```



Observations: To vary the number of ID/boarding-pass check stations and personal scanning stations, I iterated over combinations of different numbers of both types of stations. The range for each station type is between 10 and 40. Among 961 combinations, there are 185 combinations that result in an average travelers' waiting time of 15 minutes or less. The optimal number of each station can be determined based on external factors such as operation cost. However, if we assume the costs are the same, then the optimal setup includes 27 ID stations and 28 scanning stations, resulting in an average travelers' waiting time of approximately 13 minutes.

num_id_stations	num_scan_stations	avg_waiting	config
522	26	36	14.988471 26&36
525	26	39	14.979172 26&39
545	27	28	13.235261 27&28
548	27	31	14.326190 27&31
549	27	32	13.821256 27&32
...
956	40	36	4.603181 40&36
957	40	37	3.373735 40&37
958	40	38	2.178493 40&38
959	40	39	1.693107 40&39
960	40	40	1.837002 40&40

185 rows x 4 columns

```

1 optimal["stations_total"] = optimal["num_id_stations"] + optimal["num_scan_stations"]
2 optimal.loc[optimal["avg_waiting"] <= 15].sort_values(by=["stations_total"], ascending=True).head(5)

num_id_stations  num_scan_stations  avg_waiting  config  stations_total
545              27              28  13.235261  27&28             55
606              29              27  13.434151  29&27             56
576              28              28  14.727944  28&28             56
607              29              28  14.784725  29&28             57
577              28              29  12.306962  28&29             57

```